

Quantum Neural Architecture

Formal Specification of Hybrid Classical-Quantum Distributed Ingestion & Vector Processing Pipelines

V2.4 PRODUCTION SPECIFICATION

Agenda & Table of Contents

Presentation Outline

2 Agenda & Table of Contents

3 System Topology & Neural Dispatch

4 Empirical Performance Evaluation

5 Comparative Analysis & Mathematical Model


6 Cluster Analytics & Telemetry

7 Implementation & Runtime Kernel

System Topology & Neural Dispatch

Vector-Compiled Distributed Ingestion Mesh

Multi-Tier Distributed Neural Processing Pipeline



Empirical Performance Evaluation

Measured Benchmarks Across Compute Clusters

Inference Latency

BENCHMARKED SPEEDUP

1.42ms

End-to-end latency reduction under heavy concurrent load.

Throughput Capacity

BENCHMARKED SPEEDUP

450k op/s

Peak distributed operations sustained without buffer degradation.

Comparative Analysis & Mathematical Model

Formal Energy Convergence Formulation

Legacy Classical CPU / GPU Pipeline

LEGACY INFRASTRUCTURE

Linear memory scaling with quadratic communication overhead during multi-node synchronization.

AVG CONVERGENCE TIME

68ms

vs

Yumia Neural Tensor Pipeline

DISTRIBUTED TENSOR CORE

Logarithmic latency scaling leveraging quantum sparse tensor compression.

AVG CONVERGENCE TIME


4.1ms

$$\mathcal{L}(\theta) = \mathbb{E}_{x \sim \mathcal{D}} \left[\| f_{\theta}(x) - y \|_2^2 \right] + \lambda \sum_{i=1}^N \Omega(W_i)$$


Cluster Analytics & Telemetry

Resource Allocation & Workload Distribution

Inference Acceleration (x Baseline)



Cluster Resource Utilization



Implementation & Runtime Kernel

Zero-Trust Mesh & Compiler Directives

```
1 import { YumiaCompiler, createEngine } from '@yumiamd/core';  
2 const engine = createEngine({ target: 'multi-target' });  
3 const result = await engine.compile('presentation.yumia');
```

Production Verified

All cluster nodes reporting optimal throughput with zero memory leaks.

ZERO-TRUST MTLS

WCAG AAA VERIFIED